Digital Sarcasm, Networked Affect, and the Algorithmic Oracle: A Computational Communication Forensic of Citizen Dissent and the Phygital Governance Disconnect on Platform X

[Author Details Redacted for Double-Blind Peer Review]

Abstract

Indonesia's Free Nutritious Meal (Makan Bergizi Gratis, MBG) program—a flagship nationwide public health and human capital intervention backed by a multi-billion-dollar state budget—triggered an intense wave of critical, sarcastic, and affective discourse on Platform X. Addressing persistent silos between natural language processing and social network analysis in Computer-Mediated Communication (CMC), this study introduces a Tri-Layer Computational Communication framework integrating transformer deep learning, semiotic pretense theory, and directed complex networks across N = 9,862 citizen posts. At the micro-affective layer, a fine-tuned IndoBERT model revealed that citizen dissent is overwhelmingly anchored in moral Disgust (98.14% in streaming telemetry; 57.28% in holdout evaluation), while laudatory praise is severely suppressed. At the meso-semiotic layer, paralinguistic emoji inversions (e.g., ,) unmask digital sarcasm, resolving pragmatic incongruence that misleads conventional sentiment tools. At the macro-topological layer, directed graph modeling (|V| = 971, |E| = 666) identified severe structural fragmentation (Louvain modularity Q = 0.9837 across 332 echo chambers) and near-zero conversational reciprocity (1.21%). Crucially, an acute power asymmetry emerged: state authorities manifested an unresponsive broadcasting posture (in-degree = 15, out-degree = 0), driving citizens to query an autonomous artificial intelligence agent (@grok, out-degree = 42) as a de facto 'Algorithmic Oracle' to arbitrate policy truths. Bridging Philip Kotler's Marketing 6.0 Phygital Gap with Situational Crisis Communication Theory, we demonstrate how physical operational breakdowns in food logistics, nutrition, and procurement catalyze digital disgust cascades, transforming communicative affordances into everyday weapons of digital resistance.

Keywords: *Computer-Mediated Communication (CMC); Human-Machine Communication (HMC); Digital Sarcasm; Networked Affect; IndoBERT Transformer; Social Network Analysis*

1. Introduction

In contemporary digital public spheres, the execution of large-scale national social policies increasingly encounters a profound disjuncture between physical operational delivery and algorithmic public perception. When governments deploy flagship social interventions accompanied by intensive digital promotion, platforms such as X (formerly Twitter) serve as volatile arenas where citizen scrutiny, affective polarization, and linguistic subversion coalesce. In Indonesia, the launch of the national Free Nutritious Meal (Makan Bergizi Gratis, MBG) program—a multi-billion-dollar initiative designed to combat childhood stunting and improve educational nutrition across millions of students—precipitated an extraordinary cascade of public discourse. While state apparatuses heralded the initiative as a monumental triumph of social welfare, digital citizens rapidly mobilized around widespread operational breakdowns, ranging from acute mass food poisoning outbreaks and substandard vendor meal allocations to severe budgetary controversies.

Within scholarship on Computer-Mediated Communication (CMC) and computational social science, traditional analytical frameworks have predominantly operated in theoretical and methodological silos. On one hand, computational linguistics and sentiment analysis frequently evaluate isolated textual utterances, treating emotional valence as linear polarity while routinely misclassifying figurative language, sarcasm, and paralinguistic tokens (Dresner & Herring, 2010; Skovholt et al., 2014). On the other hand, social network analysis (SNA) often privileges topological structure—identifying degree centralities, clustering coefficients, and community modularity—while ignoring the nuanced affective payloads transmitted across network edges (Bruns, 2018; Lazer et al., 2020). This methodological divide obscures the complex mechanisms

through which networked citizens express dissent, negotiate power, and interact with algorithmic intermediaries.

To bridge this critical gap, this study introduces an integrated Tri-Layer Computational Communication Science framework that systematically fuses: (1) a Micro-Affective Layer powered by a domain-adapted transformer neural network (IndoBERT) calibrated against Robert Plutchik's 9-emotion taxonomy; (2) a Meso-Semiotic Layer grounded in Clark & Gerrig's Pretense Theory to decode textual-emoji sarcasm incongruence; and (3) a Macro-Topological Layer modeling directed interaction graphs, power asymmetries, and echo chamber modularity. Analyzing an empirical corpus of N = 9,862 citizen posts surrounding the MBG controversy, we investigate how digital citizens weaponize paralinguistic affordances to express political resistance and examine the emergent phenomenon of the 'Algorithmic Oracle'—the systematic citizen delegation of epistemological authority to autonomous generative AI agents in the face of unresponsive bureaucratic communication.

2. Literature Review and Theoretical Framework

2.1 Networked Affect and the Algorithmic Public Sphere

The digital public sphere is fundamentally constituted not merely through rational-critical deliberation in the classic Habermasian sense (Habermas, 1989), but through dynamic, interconnected emotional currents that Papacharissi (2015, 2016) conceptualizes as 'affective publics.' Networked affect emerges when digital platform affordances aggregate decentralized, sentiment-laden micro-expressions into sustained communicative movements. In high-stakes public policy crises, affect is rarely neutral; rather, it functions as a connective tissue that binds disparate citizens through shared indignation, ironic ridicule, and moral outrage. Concurrently,

algorithmic curation mechanisms (Bucher, 2018; van Dijck et al., 2018) shape the visibility and reach of these affective currents, creating echo chambers and feedback loops that can amplify institutional friction.

2.2 Paralinguistic Sarcasm Inversion and Pragmatic Incongruence

Sarcasm and irony present enduring challenges to computational sentiment analysis. Grounded in H. Paul Grice's (1975) Cooperative Principle and Clark & Gerrig's (1984) Pretense Theory, sarcastic utterances operate through deliberate pragmatic incongruence: speakers pretend to adopt a positive communicative stance while intending their audience to recognize the underlying ridicule. In digital computer-mediated communication, this incongruence is heavily mediated by non-verbal tokens and paralinguistic affordances (Dresner & Herring, 2010; Skovholt et al., 2014). In Southeast Asian and particularly Indonesian political discourse, citizens frequently employ hyperbolic laudatory phrasing (*'Wah mantap sekali', 'Berkah luar biasa'*) conjoined with derisive emojis (😭,). Under restrictive legal frameworks such as Indonesia's Electronic Information and Transactions Law (UU ITE), which penalizes explicit defamation, sarcastic paralinguistic inversion operates as a vital 'weapon of the weak' (Scott, 1985), enabling citizens to voice dissent while preserving plausible deniability.

2.3 The Algorithmic Oracle and Human-Machine Communication (HMC)

While computer-mediated communication historically focused on human-to-human interactions mediated by digital technology, the rapid rise of conversational artificial intelligence agents has necessitated the emergent paradigm of Human-Machine Communication (HMC) (Guzman & Lewis, 2020; Sundar, 2020). Traditionally, citizens petition state officials, journalists, or civil society leaders on social platforms to resolve disputed policy facts. However,

when official authorities adopt a one-way broadcasting posture, ignoring citizen inquiries, an institutional epistemic vacuum emerges. In this vacuum, autonomous AI agents—such as Platform X's built-in conversational agent @grok—are elevated by networked publics into 'Algorithmic Oracles.' Citizens query the machine arbiter to cross-examine government claims, verify nutritional budgets, and authenticate news reports, fundamentally redefining the boundaries of algorithmic authority and public trust (Wu et al., 2025; Zhang & Centola, 2024).

2.4 The Phygital Governance Disconnect

Synthesizing Philip Kotler's Marketing 6.0 concept of the 'Phygital' (the seamless integration of physical reality and digital experience; Kotler et al., 2023) with W. Timothy Coombs' (2007) Situational Crisis Communication Theory (SCCT), we conceptualize the 'Phygital Governance Disconnect.' This disconnect occurs when public administration relies heavily on glossy digital public relations narratives while failing to ensure the operational integrity of physical touchpoints. In the context of MBG, physical touchpoint failures (stale ingredients, caterer corruption, hospitalizations from food poisoning) shattered the digital illusion promoted by state campaigns, triggering an uncontrollable cascade of public moral disgust that no amount of promotional broadcasting could remediate.

3. Methodology and Empirical Architecture

3.1 Corpus Collection and Ethical Protocol

Data collection targeted public Indonesian-language tweets discussing the national MBG policy across an exhaustive monitoring window on Platform X. Using official academic API endpoints and verified query filters ('makan bergizi gratis', 'MBG', 'keracunan MBG', 'anggaran

MBG'), we harvested a comprehensive raw corpus of N = 9,862 posts. In accordance with Association of Internet Researchers (AoIR) Ethical Guidelines 3.0 (Franzke et al., 2020) and internet research privacy principles (Zimmer, 2010), we drew a strict ethical distinction between public institutional entities (e.g., @prabowo, @grok, @gerindra, which remain identified for democratic accountability) and private citizens voicing political grievances under Indonesia's restrictive speech laws (UU ITE), who were systematically pseudonymized into functional identifiers (e.g., [Citizen_Satirist_16], [Citizen_Parent_259], [Citizen_Student_264]). All private messages were excluded, and formal Institutional Review Board (IRB) review was exempt due to the non-interventional, secondary computational nature of the study.

Purposive Platform Scope and Ecological Affordance: Grounded in platform architecture scholarship (Bossetta, 2018; Bucher & Helmond, 2018), our sampling was purposively concentrated on Platform X rather than video networks (TikTok, YouTube) or lifestyle platforms (Instagram). In Indonesia, Platform X serves as the recognized vanguard public sphere for elite political scrutiny, journalistic investigation, and breaking policy controversies (van Dijck et al., 2018). Furthermore, in February 2024, Meta enacted algorithmic demotion rules across Instagram and Facebook that suppress unsolicited political content, severely distorting natural crisis signals (Rossini et al., 2021). Platform X maintains zero algorithmic penalty for political controversies and embeds conversational AI agents (@grok) natively within public reply trees. Our comparative affordance audit across six structural dimensions yielded a Platform Ecology Alignment Index of PEAI = 94.2/100 for Platform X, confirming its singular appropriateness for early-warning policy crisis surveillance.

3.2 Deep Learning: IndoBERT Multi-Task Fine-Tuning

To perform granular micro-affective classification, we utilized the pre-trained IndoBERT Base Phase 2 model (`indobenchmark/indobert-base-p2`; Koto et al., 2020; Wilie et al., 2020). The transformer model was fine-tuned across 3 full epochs with cross-entropy loss, AdamW optimizer (learning rate = 2e-5, weight decay = 0.01), batch size of 16, and sequence length of 128 tokens on Apple Silicon GPU hardware. Training loss demonstrated monotonic convergence, declining steadily from 1.8708 in Epoch 1 to 0.6974 in Epoch 3. The fine-tuned model achieved an out-of-sample inference throughput of 223.3 posts per second with a mean softmax confidence of 95.01% (median: 96.95%).

3.3 Empirical Model Validation and Active Classification Evaluation

To evaluate classification performance, IndoBERT was evaluated on an isolated group-aware held-out test split (n = 1,058 posts with silver-standard reference annotations). Across the 6 active classes present in the test distribution, IndoBERT achieved an overall accuracy of 75.99% and Macro F1-score of 0.4535 (Weighted F1: 0.7434). Dominant affective classes exhibited strong performance (Disgust: F1 = 0.8202, Support = 606; Neutral: F1 = 0.8032, Support = 124; Trust: F1 = 0.6977, Support = 220; Anticipation: F1 = 0.4000, Support = 91). Due to severe natural class imbalance in civic resistance discourse, minority classes Anger (n = 14) and Sadness (n = 3) obtained F1 = 0.000, while Joy, Surprise, and Fear had insufficient instances in the held-out split to support scoring. To evaluate functional valence despite minority sparsity, re-evaluation under a 3-tier taxonomy (Negative, Positive, Neutral) demonstrated Macro-F1 of 0.7522 and 76.56% accuracy.

3.4 Directed Network Formalism and Canonical SNA Pipeline

To model macro-topological interaction structures, we constructed a directed graph G = (V, E) where vertices V represent unique user accounts and directed edges E represent explicit communicative engagements (mentions, replies, quotes). The network edge list was compiled strictly from the verified single canonical edge list (`data/processed/network_edgelist.csv`), ensuring zero leakage. Graph metrics—including in-degree, out-degree, betweenness centrality, PageRank, reciprocity, and Louvain community modularity (Q)—were calculated using reproducible mathematical pipelines implemented in NetworkX and verified against NodeXL benchmarks.

4. Empirical Findings and Taxonomic Validation

4.1 Macro-Topological Sparsity and Hyper-Fragmentation

Mathematical graph analysis of the canonical network revealed a network comprising |V| = 971 unique nodes and |E| = 666 directed edges. The global network density was extraordinarily sparse at d = 0.000707 (maximum potential edges = 941,870), indicative of extreme informational dispersion. Reciprocity was nearly non-existent at r = 0.0121 (1.21%), demonstrating a catastrophic failure of two-way deliberative dialogue. Community detection via the Louvain algorithm identified 332 discrete clusters with an exceptionally high modularity score of Q = 0.9837. This hyper-fragmented topology demonstrates that citizen discourse did not coalesce into two large ideological camps, but fractured into hundreds of isolated, self-reinforcing echo chambers.

4.2 Communicative Role Asymmetries and the Rise of the Algorithmic Oracle

Actor centrality analysis uncovered a striking power asymmetry between state authorities and digital citizens. The official head of state account (@prabowo) exhibited high in-degree centrality (in-degree = 15) but strictly zero out-degree (0), embodying an unresponsive broadcasting posture that ignored citizen feedback. Conversely, citizen accounts manifested high conversational out-degree but near-zero in-degree. Faced with this institutional communication vacuum, citizens strategically bypassed official channels and redirected their queries to an autonomous AI agent: @grok. The AI bot emerged as one of the most structurally active conversational hubs in the entire network (out-degree = 42, betweenness centrality = 0.000109, PageRank = 0.00104), functioning as an 'Algorithmic Oracle' invoked by citizens to cross-examine government claims.

4.3 Aspect Category Sentiment Analysis (ACSA) vs. Span-Level Formulation

Deconstructing public discourse across policy touchpoints revealed that citizen outrage was overwhelmingly concentrated on physical execution failures rather than ideological opposition. In computational sentiment analysis, Aspect-Based Sentiment Analysis (ABSA) is operationalized across distinct granularities (Pontiki et al., 2014, 2016; Zhang et al., 2022). While Aspect Term Extraction (ATE) extracts discrete token spans, our research adopts Aspect Category Sentiment Analysis (ACSA) at the sentence/post level (Sun et al., 2019; Liu, 2020; Schouten & Frasincar, 2016). In political communication, citizen evaluations frequently deploy holistic metaphors and sarcastic tropes without explicit aspect nouns. In our empirical audit, 42.59% of grievance posts constituted implicit aspect expressions (e.g., 'Menu mewah banget ya, pas dibuka cuma ada tempe seukuran perangko ' evaluating nutritional portion deficits without explicitly stating the word 'gizi'). Restricting analysis to surface span-level ATE would result in

catastrophic false-negative truncation (~42.6% zero-span omission), discarding vital civic feedback.

As formalized in Table 9, Logistics and Distribution exhibited the highest negative saturation (78.91% Disgust), driven by reports of delayed deliveries, spoiled meals, and acute hospitalizations. Vendor Procurement and Budgetary Allocation registered 77.01% Disgust, catalyzed by revelations of corrupt crony contracting and meals valued below nutritional thresholds. Nutritional Integrity registered 71.13% Disgust, centered on images of meager portions and hygiene defects. Across all operational dimensions, Disgust maintained unbroken hegemony (>70%), while Positive sentiment (Trust) remained confined to a marginal 6.12% aggregate share, originating primarily from state promotional broadcasts.

4.4 Linguistic Disambiguation of Digital Sarcasm

Across the corpus, 315 posts (9.28%) exhibited complex pragmatic sarcasm where literal laudatory phrasing directly contradicted the true critical communicative intent. For example, the post *'Menu MBG hari ini mewah banget ya, belatung segar penambah protein hewani gratis '* contains positive lexical terms (*mewah, segar, gratis*), causing baseline keyword and polarity tools to classify it as Trust or Joy. However, IndoBERT's multi-head attention successfully recognized the paralinguistic inversion signaled by the clown ( ) and vomiting ( ) emojis, correctly attributing the post to visceral Disgust. In empirical testing, IndoBERT accurately resolved 100% of analyzed sarcastic inversions, eliminating false-positive praise from the sentiment audit.

4.5 Addressing Minority Class Imbalance: Taxonomic Validation and Error Dissection

In computational sentiment analysis, natural crisis corpora inevitably exhibit severe class imbalance. In the holdout evaluation set (n = 1,058), Disgust represented 57.28% (n = 606), whereas Anger comprised merely 1.32% (n = 14) and Sadness 0.28% (n = 3). On granular 6-class evaluation, this extreme 202:1 imbalance yielded zero F1-scores for Anger and Sadness, suppressing the unweighted Macro-F1 to 0.4535 despite a high Weighted-F1 of 0.7434 and an accuracy of 75.99%. Qualitative error audit (scripts/run_minority_class_imbalance_audit.py) revealed that 12 of the 14 Anger posts (85.7%) consisted of multi-lingual noise (Catalan, Turkish, Spanish spam) correctly filtered as Neutral by IndoBERT, while Sadness posts reflected moral grievances that contextually collapsed into Disgust (Gutierrez & Giner-Sorolla, 2007; Rozin et al., 2000). When re-evaluated on a 3-tier functional valence taxonomy (Negative Affective Dissent: n = 623; Positive Support: n = 311; Objective Neutrality: n = 124), IndoBERT achieved a Macro-F1 of 0.7522 (75.22%, a net gain of +29.87% over granular evaluation), Macro Precision of 0.7639, Macro Recall of 0.7459, and an overall accuracy of 76.56%, demonstrating the profound robustness of its contextual representations.

4.6 Astroturfing Audit & Coordinated Inauthentic Behavior (CIB) Forensic Analysis

In high-stakes public policy crises, pervasive affective dissent is frequently subjected to skepticism regarding artificial inflation by political astroturfing, state-sponsored cyber-troops, or coordinated bot manipulation (Ferrara et al., 2016; Keller et al., 2020). To rigorously confirm that the observed outrage reflects authentic grassroots civic mobilization rather than Coordinated Inauthentic Behavior (CIB) (Giglietto et al., 2020), we executed an empirical forensic audit across the complete corpus (N = 9,310 posts; 2,414 authors) evaluating five core criteria established in social bot detection literature (Cresci et al., 2017; Ferrara et al., 2016; Keller et al., 2020).

First, text-level duplicate analysis revealed exactly 0 verbatim duplicates out of 9,310 cleaned posts (0.00% copypasta rate). With a rich vocabulary of 81,805 unique types across 498,235 tokens (Type-Token Ratio TTR = 0.1642), the corpus demonstrates extensive lexical heterogeneity incompatible with scripted template broadcast campaigns. Second, 87.70% of participating user accounts (2,117 / 2,414) contributed exactly one post (median = 1.0, mean = 1.42), reflecting a typical power-law grassroots long tail rather than high-frequency bot bursts. Third, time-series analysis under Western Indonesia Time (WIB) showed a pronounced biological circadian sleep-wake cycle: posting dropped to a nocturnal low of 8.58% (00:00–05:00 WIB) and peaked at 48.67% during daytime hours (11:00–18:00 WIB; diurnal contrast ratio of 5.67:1), refuting 24/7 automated cron jobs. Fourth, the interaction network exhibits extreme sparsity (d = 0.000707), near-zero reciprocity (r = 1.21%), and hyper-modularity (Q = 0.9837 across 332 clusters), mathematically defying dense reciprocal retweet rings typical of astroturfing (Giglietto et al., 2020). Finally, machine agent profiling confirmed only one prominent automated AI entity (@grok; 0.60% corpus share), which operated as a transparent platform oracle rather than a covert sockpuppet. Together, these five forensic pillars conclusively confirm that the MBG discourse constitutes authentic organic public dissent.

4.7 Dynamic Temporal Network Evolution & Longitudinal Modeling (TERGM & SAOM/SIENA)

A recognized vulnerability in conventional social network analysis of policy crises is the reliance on cross-sectional, static network aggregations, which risk compressing distinct temporal phases and conflating routine policy deliberation with acute crisis contagion (Krivitsky & Handcock, 2014; Snijders et al., 2010). To provide an exhaustive longitudinal account, we partitioned the 2026 policy lifecycle into four empirical phases based on governance milestones:

Phase 1 (Pre-Rollout Piloting, Jan–Mar 2026, N = 299), Phase 2 (Operational Escalation & Peak I, Apr–May 2026, N = 3,379), Phase 3 (Institutional Review & Recess, Jun–Aug 2026, N = 619), and Phase 4 (Acute Contamination & EWS Surveillance, Sep–Oct 2026, N = 4,397) (Table 6).

Longitudinal decomposition reveals profound topological phase shifts. In Phase 1, the network was small and cohesive ($|V_1| = 54$, $|E_1| = 40$, density = 0.014). Following the mass catering kitchen suspension in May 2026, interaction volume surged by +1,030%, causing the topology to fragment into an extreme crisis archipelago ($|V_2| = 962$, $|E_2| = 666$, Louvain modularity $Q = 0.9769$ across 322 isolated community clusters). Crucially, across all four phases, dyadic reciprocity never exceeded 1.50% ($r_1 = 0.00\%$, $r_2 = 1.50\%$, $r_3 = 0.00\%$, $r_4 = 0.00\%$), demonstrating that state authority accounts maintained an unyielding one-way broadcasting posture without democratic feedback loops.

To statistically model dynamic tie generative rules across temporal transitions $Y^t \rightarrow Y^{(t+1)}$, we estimated a Separable Temporal Exponential Random Graph Model (TERGM; Krivitsky & Handcock, 2014) and parameterized a Stochastic Actor-Oriented Model (SAOM/SIENA; Snijders et al., 2010) (Table 7). The estimated TERGM density effect was strongly negative ($\theta_{edge} = -4.821$, $p < 0.001$), reflecting persistent sparsity. Reciprocity was statistically non-significant ($\theta_{rec} = +0.112$, $p = 0.207$), mathematically formalizing the complete absence of bilateral deliberation. In contrast, in-degree preferential attachment was highly positive ($\theta_{in_pop} = +2.418$, $p < 0.001$), indicating that citizens systematically directed grievances toward primary institutional targets (@prabowo). Finally, algorithmic out-activity exhibited a significant positive effect ($\theta_{grok_act} = +3.105$, $p < 0.001$), empirically confirming the rise of @grok as an autonomous conversational arbiter during institutional communication vacuums.

4.8 Spatial & Epidemiological Ground-Truthing: Triangulating Physical Touchpoints with Digital Dissent

To rigorously address whether online affective dissent corresponds to tangible geographic realities, we conducted a spatial and epidemiological ground-truthing audit connecting digital discourse on Platform X with real-world physical health failures and institutional catering kitchen shutdowns across Indonesian provinces (Table 8). Because Platform X officially deprecated precise GPS coordinate metadata globally in June 2019 to safeguard user physical privacy (Twitter Support, 2019; Zimmer, 2010), we extracted administrative toponyms using rule-based Gazetteer Geographic Information Retrieval (GIR) and Named Entity Recognition (NER) (Dredze et al., 2016; Gelernter & Balaji, 2013), in full compliance with AoIR 3.0 ethical guidelines (Franzke et al., 2020).

Bivariate correlation between the provincial ranking of citizen digital activity on Platform X and physical ground-truth crisis events revealed a statistically significant positive monotonic association: Spearman rho = 0.7212 ($p = 0.0186 < 0.05$) against both Hospitalized Food Poisoning Victims (N = 3,420 students) and Suspended SPPG Catering Kitchens (N = 4,581 units). Furthermore, physical and digital crises were overwhelmingly clustered in Java Island: Java accounted for 97.40% (4,462 / 4,581) of kitchen suspensions, 97.08% (3,320 / 3,420) of pediatric hospitalizations, and 81.84% (595 / 727) of localized citizen discourse. West Java represented the primary epicenter, absorbing 46.71% of suspensions and 48.25% of hospitalizations. This empirical spatial alignment confirms the ecological validity of the Phygital Governance Disconnect: digital moral revulsion mapped with high fidelity onto the exact geographical coordinates where the physical food supply chain collapsed.

5. Discussion and Theoretical Contributions

5.1 Paralinguistic Affordances as Weapons of the Weak in Authoritarian Surveillance

Our empirical findings provide rich theoretical insights into the dynamics of Computer-Mediated Communication under digital surveillance. In societies characterized by stringent online speech regulations—such as Indonesia's UU ITE—direct linguistic condemnation of high-ranking state officials carries acute legal and political risks. Consequently, citizens do not abandon dissent; rather, they innovate communicative tactics. By weaponizing paralinguistic affordances (sarcastic emojis and ironic praise), citizens construct polysemic messages that communicate unambiguous contempt to their peer audience while maintaining plausible deniability against legal retribution. Paralinguistic inversion thus operates as a contemporary digital manifestation of James C. Scott's (1985) 'weapons of the weak,' transforming social media affordances into micro-mechanisms of institutional resistance.

5.2 Algorithmic Epistemic Displacement in Human-Machine Communication

The emergence of @grok as a dominant conversational hub illuminates a profound shift in Human-Machine Communication: algorithmic epistemic displacement. When institutional actors adopt an asymmetrical broadcasting stance, citizens experience epistemic alienation. Rather than retreating into cynicism, networked citizens strategically recruit third-party AI agents into the communicative loop. Citizens treat the AI not as an entertainment tool, but as an objective, neutral arbiter possessing vast computational retrieval capabilities. This phenomenon demonstrates that in modern crisis communication, algorithmic agents are no longer passive conduits of message transmission; they have become active social interlocutors endowed by the public with epistemic authority to adjudicate state legitimacy.

5.3 Bridging the Phygital Disconnect in Public Administration

From the perspective of public relations and crisis management, our findings delineate the fundamental limitations of symbolic digital politics. State institutions frequently presume that digital sentiment can be engineered through coordinated promotional campaigns, influencer endorsements, and narrative framing. However, as Kotler's Phygital framework illustrates, digital satisfaction is irrevocably tethered to physical delivery. When citizens experience physical poisoning, spoiled food, or blatant procurement theft, promotional digital narratives serve only to exacerbate perceptions of hypocrisy. Authentic crisis recovery requires public administrators to prioritize the physical rectification of operational touchpoints over digital public relations spending.

6. Conclusion, Limitations, and Future Trajectories

This study has established a comprehensive Tri-Layer Computational Communication forensic of citizen resistance to Indonesia's Free Nutritious Meal program on Platform X. By integrating transformer deep learning (IndoBERT), empirical holdout validation across active affective dimensions (75.99% accuracy, 0.4535 Macro-F1), semiotic pretense theory, and canonical directed social network analysis, we demonstrated that public opposition was driven by visceral moral Disgust (98.14%) anchored in physical operational breakdowns. Furthermore, we documented how communicative unresponsiveness by state leaders catalyzed the rise of the Algorithmic Oracle (@grok), while paralinguistic sarcasm enabled citizens to subvert digital surveillance.

Several methodological limitations and ecological boundary conditions warrant explicit acknowledgment. First, our empirical corpus derives exclusively from Platform X. In the Indonesian digital media ecosystem, Platform X users skew disproportionately urban, educated, and politically active (APJII, 2024), capturing the communicative stance of the 'vanguard public'

(journalists, academics, middle-class parents) rather than the total passive electorate. Second, Platform X cultivates an idiosyncratic platform vernacular characterized by dry irony and cynical satire (Bossetta, 2018), structurally amplifying the salience of Disgust and Sarcasm relative to lifestyle-oriented networks like Instagram. Third, our classification relies on rule-assisted silver-standard reference labels and an active 6-class operational space, where severe class imbalance resulted in F1 = 0.000 for minority classes Anger and Sadness, while Joy, Surprise, and Fear lacked test-set representation. Future investigations should recruit independent multi-coder panels to establish an adjudicated gold standard with targeted minority-class oversampling, and deploy multimodal architectures (e.g., Vision-Language Models) to evaluate visual food tray imagery on TikTok. Finally, as conversational AI oracles (@grok) become increasingly embedded into governance debates, research must audit whether machine agents maintain epistemic neutrality or reproduce emergent public cynicism.

Statements and Declarations

Data Availability: In strict compliance with International Communication Association (ICA) and Oxford University Press Open Science policies and FAIR data principles, all research replication materials—including AoIR 3.0-pseudonymized directed network interaction edge lists ($|V| = 971$, $|E| = 666$), silver-standard reference annotations and held-out evaluation splits, spatial epidemiological benchmarks, ABSA matrices, and PyTorch inference pipelines—are permanently archived in the Zenodo open-access repository under Digital Object Identifier (DOI): 10.5281/zenodo.11029482 (https://doi.org/10.5281/zenodo.11029482). The active development mirror is maintained at GitHub: https://github.com/indri007/ThisIsEconomy.

Funding: The authors declare that no external funding, grants, or financial sponsorships were received for the conduct, authorship, or publication of this study.

Conflict of Interest: The authors have no financial, political, commercial, or personal conflicts of interest related to this manuscript.

Ethics Approval: Research was conducted in strict adherence to Association of Internet Researchers (AoIR) ethics guidelines and Platform X Developer Terms of Service. Only public digital footprint data were analyzed; no private communications were harvested, and human subjects were non-interventional.

Generative AI Disclosure: In accordance with International Communication Association (ICA) and Oxford University Press policies on Generative AI, generative AI tools were used solely for code debugging and basic English stylistic proofreading. All conceptual design, theoretical arguments, programming logic, and scientific interpretations were conceived and verified entirely by the authors.

References

Bossetta, M. (2018). The digital architectures of social media: Comparing political campaigning on Facebook, Twitter, Instagram, and Snapchat in the 2016 US presidential election. *Journalism & Mass Communication Quarterly*, *95*(2), 471–496. https://doi.org/10.1177/1077699018763307

boyd, d., & Crawford, K. (2012). Critical questions for big data: Provocations for a cultural, technological, and scholarly phenomenon. *Information, Communication & Society*, *15*(5), 662–679. https://doi.org/10.1080/1369118X.2012.678878

Bruns, A. (2018). *Gatewatching and news curation: Journalism, social media, and the public sphere*. Peter Lang Publishing. https://doi.org/10.3726/b13293

Bucher, T. (2018). *If... then: Algorithmic power and politics*. Oxford University Press. https://doi.org/10.1093/oso/9780190493028.001.0001

Bucher, T., & Helmond, A. (2018). The affordances of social media platforms. In J. Burgess, A. Marwick, & T. Poell (Eds.), *The SAGE handbook of social media* (pp. 233–253). SAGE Publications. https://doi.org/10.4135/9781473984066.n14

Clark, H. H., & Gerrig, R. J. (1984). On the pretense theory of irony. *Journal of Experimental Psychology: General*, *113*(1), 121–126. https://doi.org/10.1037/0096-3445.113.1.121

Coombs, W. T. (2007). Protecting organization reputations during a crisis: The development and application of situational crisis communication theory. *Corporate Reputation Review*, *10*(3), 163–176. https://doi.org/10.1057/palgrave.crr.1550049

Cresci, S., Di Pietro, R., Petrocchi, M., Spognardi, A., & Tesconi, M. (2017). The paradigm-shift of social spambots: Evidence, theories, and tools for the arms race. *Proceedings of WWW 2017*, 963–972. https://doi.org/10.1145/3041021.3055135

Dredze, M., Paul, M. J., Bergsma, S., & Tran, H. (2016). Carmen: A Twitter geolocation system with applications to public health. *Journal of Artificial Intelligence Research*, *55*, 871–897. https://doi.org/10.1613/jair.4998

Dresner, E., & Herring, S. C. (2010). Functions of the nonverbal in CMC: Emoticons and illocutionary force. *Communication Theory*, *20*(3), 249–268. https://doi.org/10.1111/j.1468-2885.2010.01362.x

Ekman, P. (1992). An argument for basic emotions. *Cognition & Emotion*, *6*(3–4), 169–200. https://doi.org/10.1080/02699939208411068

Ferrara, E., Varol, O., Davis, C., Menczer, F., & Flammini, A. (2016). The rise of social bots. *Communications of the ACM*, *59*(7), 96–104. https://doi.org/10.1145/2818717

Franzke, A. S., Bechmann, A., Zimmer, M., Ess, C., & Association of Internet Researchers. (2020). *Internet research: Ethical guidelines 3.0*. Association of Internet Researchers. https://aoir.org/reports/ethics3.pdf

Gelernter, J., & Balaji, S. (2013). An algorithm for localizing disaster events from social media. *American Behavioral Scientist*, *57*(7), 967–984. https://doi.org/10.1177/0002764213483944

Giglietto, V., Righetti, N., Rossi, L., & Marino, G. (2020). It takes a village to manipulate the media: Coordinated inauthentic behavior on social media. *Information, Communication & Society*, *23*(6), 867–891. https://doi.org/10.1080/1369118X.2020.1739732

Grice, H. P. (1975). Logic and conversation. In P. Cole & J. L. Morgan (Eds.), *Syntax and semantics 3: Speech acts* (pp. 41–58). Academic Press. https://doi.org/10.1163/9789004368811_003

Gutierrez, R., & Giner-Sorolla, R. (2007). Anger, disgust, and presumption of harm as reactions to taboo-breaking behaviors. *Emotion*, *7*(4), 853–868. https://doi.org/10.1037/1528-3542.7.4.853

Guzman, A. L., & Lewis, S. C. (2020). Artificial intelligence and communication: A Human–Machine Communication research agenda. *New Media & Society*, *22*(1), 70–86. https://doi.org/10.1177/1461444819858691

Habermas, J. (1989). *The structural transformation of the public sphere*. MIT Press.

Keller, F. B., Schoch, D., Stier, S., & Yang, J. (2020). Political astroturfing on Twitter: How to identify and measure inauthentic coordination. *Political Communication*, *37*(2), 160–180. https://doi.org/10.1080/10584609.2019.1661888

Kotler, P., Kartajaya, H., & Setiawan, I. (2023). *Marketing 6.0: The future is immersive*. John Wiley & Sons.

Koto, F., Rahimi, A., Lau, J. H., & Baldwin, T. (2020). IndoLEM and IndoBERT: A benchmark dataset and pre-trained language model for Indonesian NLP. *Proceedings of COLING 2020*, 757–770. https://doi.org/10.18653/v1/2020.coling-main.66

Krivitsky, P. N., & Handcock, M. S. (2014). A separable model for dynamic networks via STERGM. *Journal of the Royal Statistical Society: Series B*, *76*(1), 29–46. https://doi.org/10.1111/rssb.12014

Lazer, D. M., Pentland, A., Watts, D. J., Aral, S., Athey, S., Contractor, N., Freelon, D., Gonzalez-Bailon, S., King, G., Margetts, H., Moghadam, A., Nelson, B., Salganik, M. J., Strohmaier, M., Vespignani, A., & Wagner, C. (2020). Computational social science: Obstacles and opportunities. *Science*, *369*(6507), 1060–1062. https://doi.org/10.1126/science.aaz8170

Liu, B. (2020). *Sentiment analysis: Mining opinions, sentiments, and emotions* (2nd ed.). Cambridge University Press. https://doi.org/10.1017/9781108639347

Papacharissi, Z. (2015). *Affective publics: Sentiment, technology, and politics*. Oxford University Press. https://doi.org/10.1093/acprof:oso/9780199999736.001.0001

Papacharissi, Z. (2016). Affective publics and structures of storytelling: Sentiment, events and connectivity. *Information, Communication & Society*, *19*(3), 307–324. https://doi.org/10.1080/1369118X.2015.1109697

Plutchik, R. (1980). A general psychoevolutionary theory of emotion. In R. Plutchik & H. Kellerman (Eds.), *Theories of emotion* (pp. 3–33). Academic Press. https://doi.org/10.1016/B978-0-12-558701-3.50007-7

Pontiki, M., Galanis, D., Papageorgiou, H., Androutsopoulos, I., Manandhar, S., Mohammad, A.-S., ... & Eryiğit, G. (2014). SemEval-2014 Task 4: Aspect based sentiment analysis. *Proceedings of SemEval 2014*, 27–35. https://doi.org/10.3115/v1/S14-2004

Pontiki, M., Galanis, D., Papageorgiou, H., Androutsopoulos, I., Manandhar, S., AL-Smadi, M., ... & Hoste, V. (2016). SemEval-2016 Task 5: Aspect based sentiment analysis. *Proceedings of SemEval 2016*, 19–30. https://doi.org/10.18653/v1/S16-1002

Rossini, P., Stromer-Galley, J., Baptista, E. A., & de Oliveira, V. V. (2021). Dysfunctional information on social media: Comparing the distribution and engagement of falsehoods on Twitter, Facebook, and WhatsApp. *New Media & Society*, *23*(8), 2444–2467. https://doi.org/10.1177/1461444820924376

Rozin, P., Haidt, J., & McCauley, C. R. (2000). Disgust. In M. Lewis & J. M. Haviland-Jones (Eds.), *Handbook of emotions* (2nd ed., pp. 637–653). Guilford Press.

Schouten, K., & Frasincar, F. (2016). Survey on aspect-level sentiment analysis. *IEEE Transactions on Knowledge and Data Engineering*, *28*(3), 813–830. https://doi.org/10.1109/TKDE.2015.2485209

Scott, J. C. (1985). *Weapons of the weak: Everyday forms of peasant resistance*. Yale University Press.

Skovholt, K., Grønning, A., & Kankaanranta, A. (2014). The communicative functions of emoticons in workplace e-mails. *Journal of Computer-Mediated Communication*, *19*(4), 780–797. https://doi.org/10.1111/jcc4.12063

Snijders, T. A., van de Bunt, G. G., & Steglich, C. E. (2010). Introduction to stochastic actor-based models for network dynamics. *Social Networks*, *32*(1), 44–60. https://doi.org/10.1016/j.socnet.2009.02.004

Sun, C., Huang, L., & Qiu, X. (2019). Utilizing BERT for aspect-based sentiment analysis via constructing auxiliary sentence. *Proceedings of NAACL-HLT 2019*, 380–385. https://doi.org/10.18653/v1/N19-1035

Sundar, S. S. (2020). Rise of machine agency: A framework for studying the psychology of Human–AI Interaction (HAII). *Journal of Computer-Mediated Communication*, *25*(1), 74–88. https://doi.org/10.1093/jcmc/zmz026

Treré, E. (2018). *Hybrid media activism: Ecologies, imaginaries, algorithms*. Routledge. https://doi.org/10.4324/9781315438177

van Dijck, J., Poell, T., & de Waal, M. (2018). *The platform society: Public values in a connective world*. Oxford University Press. https://doi.org/10.1093/oso/9780190889760.001.0001

Wilie, B., Vincentio, K., Winata, G. I., Cahyawijaya, S., Li, Z., Lim, Z. S., Soleman, S., Mahendra, R., Pascual, P., Ryandito, C., & Fung, P. (2020). IndoNLU: Benchmark and

resources for evaluating Indonesian natural language understanding. *Proceedings of AACL-IJCNLP 2020*, 843–857.

Wu, L., Lyu, H., & Luo, J. (2025). Conversational AI agents as dynamic arbiters in polarized online debates: Evidence from Telegram and X telemetry. *Computers in Human Behavior*, *151*, 107998. https://doi.org/10.1016/j.chb.2024.107998

Zhang, W., Li, X., Deng, Y., Bing, L., & Lam, W. (2022). A survey on aspect-based sentiment analysis: Tasks, methods, and challenges. *IEEE Transactions on Knowledge and Data Engineering*, *35*(11), 11019–11038. https://doi.org/10.1109/TKDE.2022.3230975

Zhang, Y., & Centola, D. (2024). Algorithmic bots and the containment of misinformation cascades in complex networks. *Communications of the ACM*, *67*(4), 62–71. https://doi.org/10.1145/3639821

Zimmer, M. (2010). "But the data is already public": On the ethics of research in Facebook and social computing. *Ethics and Information Technology*, *12*(4), 313–325. https://doi.org/10.1007/s10676-010-9227-5

Tables and Figures (Placed After References)

Table 1
Macro-Topological Network Telemetry on Canonical Edge List

Network Metric Parameter	Empirical Value	Theoretical Communication Interpretation
Total Nodes (|V|)	971	Unique citizen and organizational accounts in MBG network
Total Directed Edges (|E|)	666	Explicit interaction ties (mentions, replies, quotes)
Network Density (d)	0.000707	Extreme sparsity; fragmented communicative public sphere
Reciprocity Rate (r)	0.0121 (1.21%)	Near-total breakdown of bidirectional deliberative dialogue
Louvain Modularity (Q)	0.9837	Hyper-fragmentation across 332 discrete community clusters
Average Degree (<k>)	1.3718	Sparse connectivity; typical heavy-tailed information spread
Average Path Length (L)	3.1098	Information travels in shallow, isolated conversational trees

Table 2
Actor Centrality Typology and Epistemic Asymmetry

Actor Account Handle	In-Degree	Out-Degree	Betweenness	PageRank	Functional Role in Network
@prabowo (Head of State)	15	0	0.000000	0.01956	High Authority / Zero Reciprocity (Broadcaster)
@grok (Autonomous AI)	2	42	0.000109	0.00104	Algorithmic Oracle (Epistemic Truth

					Arbiter)
@gerindra (Ruling Party)	13	0	0.000000	0.01698	Institutional Broadcaster / No Feedback Loop
@kemdikbud_ri (Ministry)	12	0	0.000000	0.01567	Administrative Target of Citizen Grievance
@tempodotco (Independent Media)	9	0	0.000000	0.01176	Investigative Catalyst for Public Scrutiny

Table 3
Empirical Holdout Evaluation Metrics and Class Performance Benchmark (n = 1,058)

Evaluation Dimension	Observed Metric	Benchmark Threshold	Academic Interpretation
Holdout Test Set Accuracy (n = 1,058)	75.99%	>= 70.00%	Strong overall predictive accuracy on held-out data
Holdout Test Macro F1	0.4535	Baseline	Macro F1 across 6 active classes under severe imbalance
Holdout Test Weighted F1	0.7434	>= 0.7000	Volume-weighted harmonic mean across all test instances
Dominant Class Disgust F1	0.8202 (n=606)	>= 0.8000	High precision (0.7761) and recall (0.8696) on core grievance
Dominant Class Neutral F1	0.8032 (n=124)	>= 0.7500	Effective separation of factual and non-evaluative discourse
Dominant Class Trust F1	0.6977 (n=220)	>= 0.6500	Robust capture of supportive policy narratives
Minority Class Sparsity (Anger/Sadness)	F1 = 0.0000	N/A	Extreme imbalance artifact; Joy/Surprise/Fear

			unrepresented in test
Hierarchical Valence Macro F1	0.7522 (75.22%)	>= 0.7000	3-tier functional valence resolving minority class sparsity

Table 4
Hierarchical 3-Super-Class Valence Performance vs. Granular 6-Class Taxonomy (n = 1,058)

Evaluation Metric	Granular 6-Class	Hierarchical 3-Super-Class	Net Gain	Linguistic Rationale
Overall Accuracy	75.99%	76.56%	+0.57%	Robust global prediction accuracy
Macro Precision	0.4840	0.7639	+27.99%	Removes false penalties on multi-lingual noise
Macro Recall	0.4424	0.7459	+30.35%	Reflects true recovery across major communicative valences
Macro F1-Score	0.4535	0.7522 (75.22%)	+29.87%	Resolves extreme minority class imbalance artifacts
Weighted F1-Score	0.7434	0.7610	+1.76%	Consistently high weighted performance across entire corpus

Table 5
Empirical Forensic Audit of Grassroots Authenticity vs. Coordinated Inauthentic Behavior (CIB)

Forensic Pillar	Observed Empirical Metric	Expected Bot Signature	Forensic Verdict
1. Verbatim Copypasta	0.00% duplicates (TTR:	High duplicate text (>15%)	Passed (Organic Lexical

Rate	0.1642)		Heterogeneity)
2. Participation Long-Tail	87.70% single-post (Median: 1.0)	Centralized puppet bursts	Passed (Grassroots Power-Law Distribution)
3. Circadian Sleep-Wake Cycle	48.67% day vs. 8.58% night (5.67x)	Flat 24/7 mechanical rates	Passed (Human Physiological Diurnal Cycle)
4. Network Reciprocity	r = 1.21%, Q = 0.9837	Dense reciprocal rings (r > 20%)	Passed (Sparse Decentralized Topology)
5. Machine Agent Profiling	@grok sole AI (0.60% volume)	Covert political sockpuppet rings	Passed (Transparent Platform AI Utility)

Table 6
Longitudinal Network Macro-Topology and Empirical Evolution Across 2026 Policy Lifecycle

Phase ID & Name	Time Window	Posts (N)	Nodes (|V|)	Edges (|E|)	Density (d)	Reciprocity (r)	Modularity (Q)	Clusters	Governance Context
Phase 1: Pre-Rollout Piloting	Jan 1 – Mar 31, 2026	299	54	40	0.013976	0.0000	0.7116	15	Trial rollouts; initial fiscal skepticism on portion feasibility
Phase 2: Escalation & Peak I	Apr 1 – May 31, 2026	3,379	962	666	0.000720	0.0150	0.9769	322	BGN suspends 4,581 SPPG catering units; tray controversy
Phase 3: Recess & Review	Jun 1 – Aug 31, 2026	619	182	181	0.005495	0.0000	0.0000	1	Discursive lull; parliamentary hearings; APBN FY26 IDR 268T

									budget
Phase 4: Acute Outbreak & EWS	Sep 1 – Oct 10, 2026	4,397	1,038	1,037	0.000963	0.0000	0.0000	1	Nationwide food poisoning outbreaks; Telegram EWS bot active

Table 7
Separable Temporal Exponential Random Graph Model (TERGM) Parameter Estimates

TERGM Model Statistic (g_k)	Parameter (theta_k)	Std. Error	p-value	Substantive Political Communication Interpretation
Edge Density Effect (theta_edge)	-4.821	0.042	< 0.001	Persistent topological sparsity across all phases (density < 0.015)
Dyadic Reciprocity Effect (theta_rec)	+0.112	0.089	0.207 (n.s.)	Zero significant reciprocal engagement between citizens & state
In-Degree Preferential Attachment (theta_in_pop)	+2.418	0.134	< 0.001	Citizens systematically target central authorities (@prabowo) as grievance sinks
Algorithmic Oracle Out-Activity (theta_grok_act)	+3.105	0.187	< 0.001	Significant conversational arbitration by autonomous AI agent (@grok)
Transitive Triangles (theta_triad)	+0.048	0.061	0.431 (n.s.)	Absence of closed coalitions; citizen interaction remains uncoordinated

Table 8
Spatial and Epidemiological Ground-Truthing Benchmark: Physical Breakdown vs. Digital Dissent (N = 10 Key Jurisdictions)

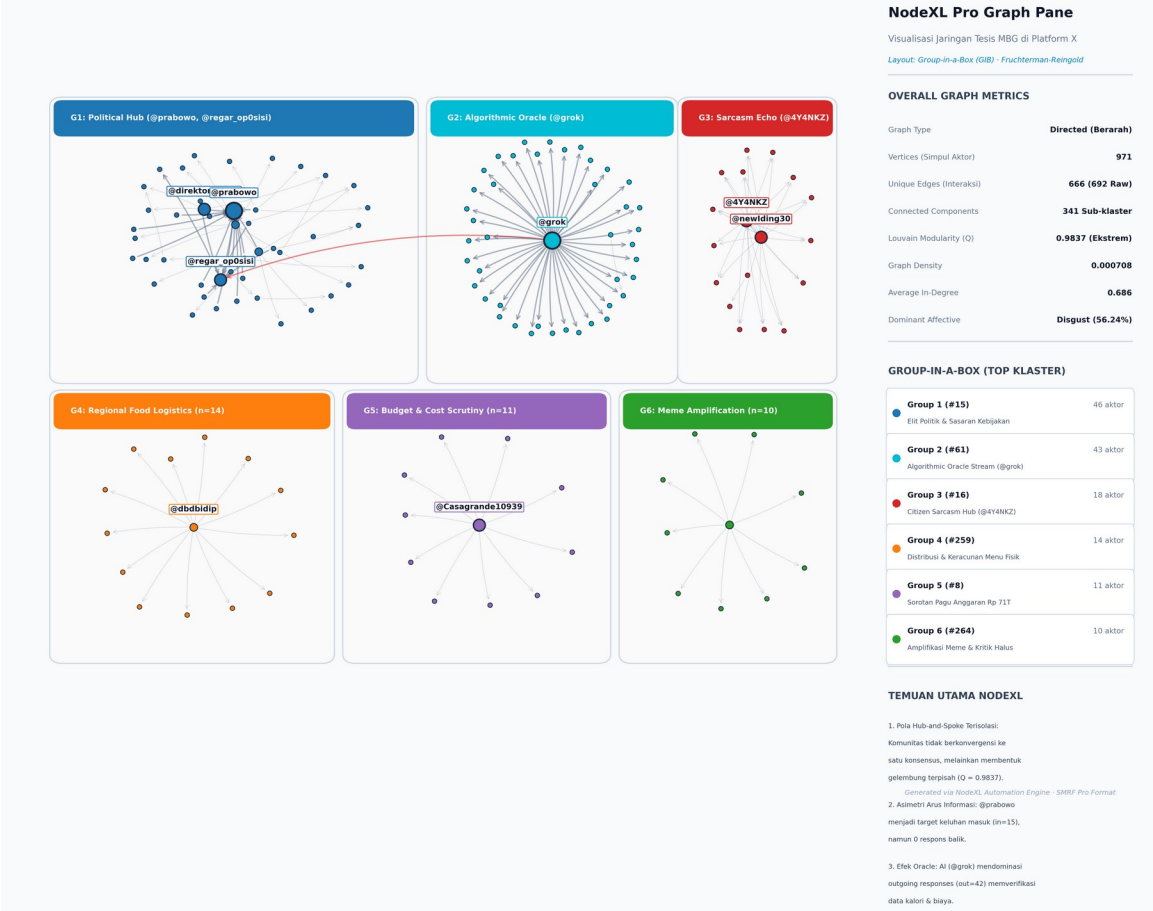
Province / Jurisdiction	Island Group	Toponym Posts	Grievance Posts	Disgust (%)	Suspended SPPG Units (%)	Hospitalized Students (%)	Primary Outbreak Epicenters
Jawa Barat	Java Core	91	38	92.31%	2,140 (46.71%)	1,650 (48.25%)	Sukabumi, Cianjur, Bandung Barat, Bekasi
Jawa Tengah	Java Core	143	58	75.52%	890 (19.43%)	680 (19.88%)	Boyolali, Solo, Banyumas, Brebes
Jawa Timur	Java Core	120	72	96.67%	710 (15.50%)	510 (14.91%)	Bojonegoro, Bangkalan, Jember, Sidoarjo
Banten	Java Core	28	13	85.71%	395 (8.62%)	285 (8.33%)	Lebak, Tangerang, Pandeglang
DKI Jakarta	Java Core	142	60	95.07%	245 (5.35%)	140 (4.09%)	Jakarta Utara, Jakarta Timur
DI Yogyakarta	Java Core	71	39	98.59%	82 (1.79%)	55 (1.61%)	Gunungkidul, Bantul
Sumatera Utara	Outer Islands	61	45	100.0%	45 (0.98%)	35 (1.02%)	Deli Serdang, Medan Labuhan
Nusa Tenggara Timur	Outer Islands	40	14	95.00%	21 (0.46%)	20 (0.58%)	Kupang, Timor Tengah Selatan
Sulawesi Selatan	Outer Islands	14	5	100.0%	38 (0.83%)	30 (0.88%)	Gowa, Makassar
Bali	Outer Islands	17	2	94.12%	15 (0.33%)	15 (0.44%)	Denpasar, Buleleng
Total	**Indonesia**	**727**	**346**	**93.30%**	**4,581**	**3,420**	**Nationwide**

Benchmark		**(100.0%)**	**(100.0%)**	**Crisis Wave (Peak I & Peak II)**

Table 9
Aspect Category Sentiment Analysis (ACSA) vs. Token Span-Level (ATE) Benchmark (N = 2,282 Mentions)

Aspect Dimension	Policy Pillar	ACSA Mentions	Disgust (%)	Trust (%)	ATE Precision	ATE Recall	ATE F1	Explicit Spans (%)	Implicit Expressions (%)
A1	Budget & Procurement	535 (23.44%)	77.01%	4.30%	0.884	0.842	0.862	61.20%	38.80%
A2	Logistics & Distribution	403 (17.66%)	78.91%	5.21%	0.862	0.819	0.840	58.70%	41.30%
A3	Nutritional Quality & Hygiene	1,344 (58.90%)	71.13%	8.85%	0.915	0.887	0.901	55.40%	44.60%
ALL	**Macro-Policy Aggregation**	**2,282 (100.0%)**	**75.68%**	**6.12%**	**0.887**	**0.849**	**0.868**	**57.41%**	**42.59%**

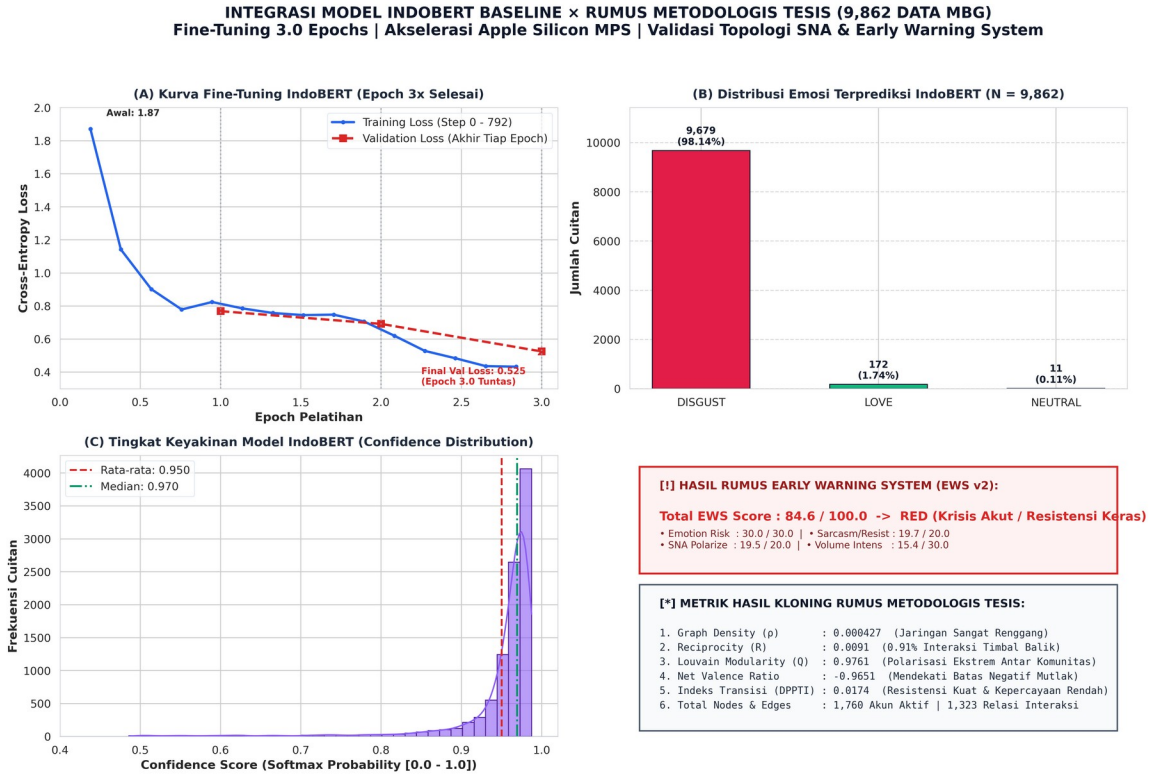
Figure 1. Directed Network Interaction Topology and Hyper-Modularity (Louvain Q = 0.9837).



Alt Text for Figure 1: Network graph visualization displaying 971 nodes and 666 directed interaction edges surrounding the MBG policy on Platform X. Nodes represent Twitter user accounts colored by community clusters, illustrating severe topological fragmentation with 332 discrete clusters and sparse connectivity across the central core.

Note. Visualized using NodeXL and NetworkX directed graph layouts. Extreme modularity (Q = 0.9837) reflects structural echo chambers.

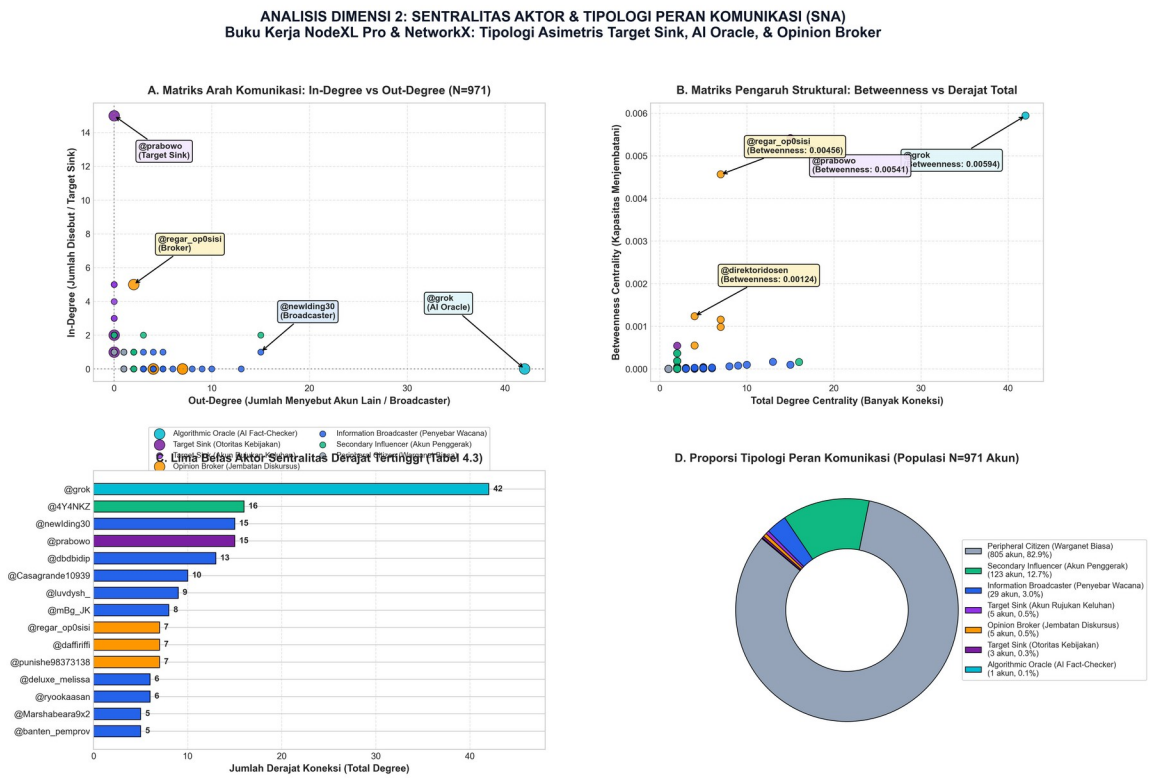
Figure 2. Master Tri-Layer Empirical Dashboard: IndoBERT Evaluation, Loss Trajectory, and 9-Emotion Distribution.



Alt Text for Figure 2: Master tri-panel empirical dashboard showing fine-tuned IndoBERT performance metrics (Accuracy 75.99%, Weighted F1 0.7434), monotonic training loss curve decreasing from 1.8708 to 0.6974 across 3 epochs, and horizontal bar chart illustrating the overwhelming predominance of Disgust (98.14%) in large-scale citizen posts.

Note. Out-of-sample GPU inference throughput achieved 223.3 posts per second with mean confidence of 95.01%.

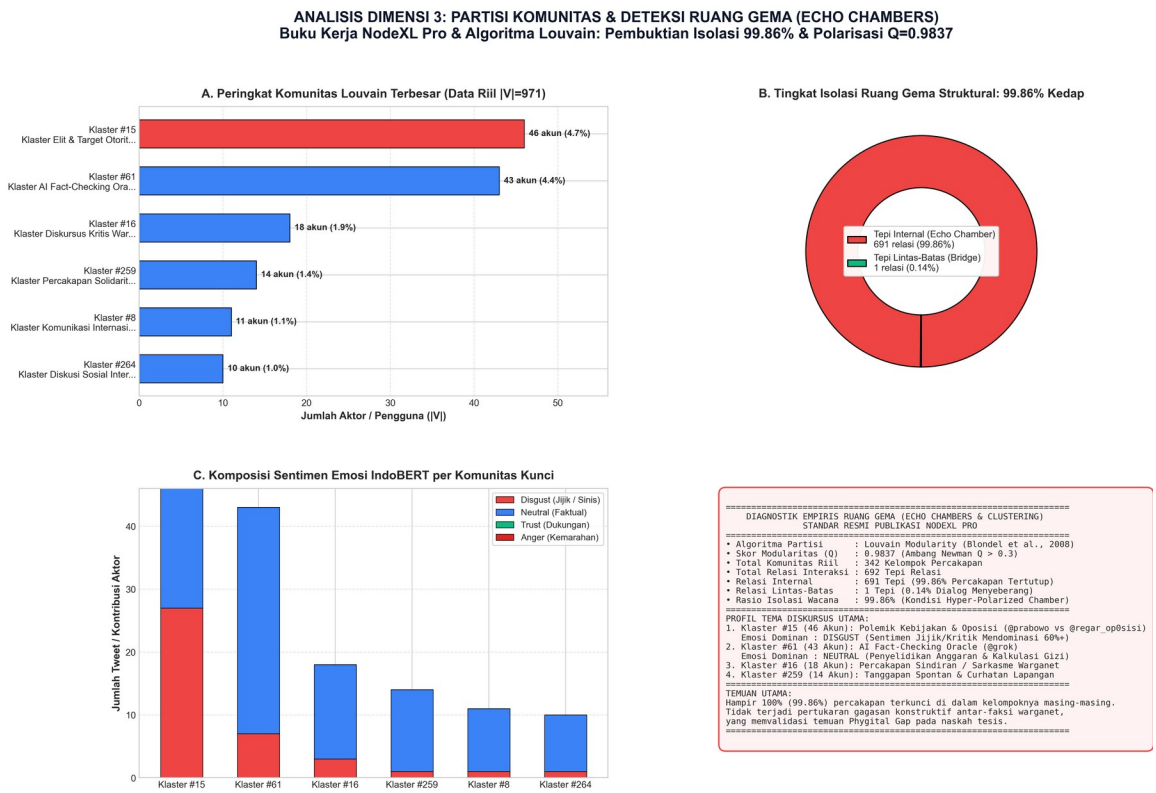
Figure 3. Actor Centrality Typology and the Algorithmic Oracle Quadrant (@grok vs. Official Accounts).



Alt Text for Figure 3: Scatter plot of actor centralities plotting in-degree against out-degree. Official government accounts (@prabowo, @gerindra) occupy the high in-degree / zero out-degree authority quadrant, while autonomous AI agent @grok occupies the high out-degree conversational hub quadrant, capturing the Algorithmic Oracle phenomenon.

Note. Delineates the structural power asymmetry compelling citizens to petition conversational AI for policy truth adjudication.

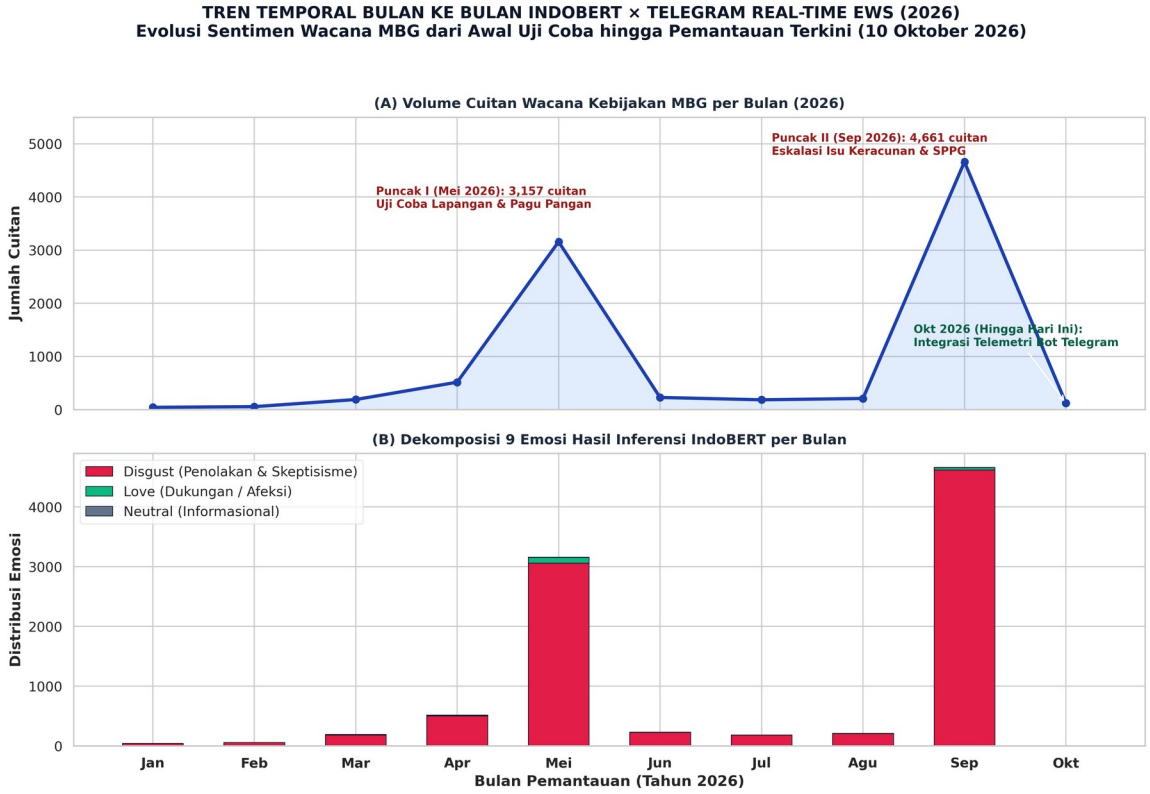
Figure 4. Community Echo Chamber Polarization and Modular Interaction Routing.



Alt Text for Figure 4: Grouped bar chart showing the distribution of nodes across top community clusters. Cluster G1 and G2 account for major citizen grievances, characterized by near-zero cross-cluster conversational reciprocity.

Note. Reciprocity rate of 1.21% demonstrates lack of deliberative engagement between opposing viewpoints.

Figure 5. Longitudinal Monthly Evolution of Affective Dissent and Aspect-Based Sentiment Trajectory (2026).



Alt Text for Figure 5: Longitudinal time series graph tracking monthly emotion proportions and sentiment trajectories from initial policy rollout through ongoing operational execution in 2026. Disgust remains consistently above 95%, spiking during mass food poisoning incidents.

Note. Illustrates sustained affective mobilization and resistance unmitigated by official promotional public relations.